

MINIMUM-INTERVENTION HUMAN CORRECTION: BRANCH-STRUCTURED EVIDENCE FOR EMBODIED WORLD-ACTION MODELS

Anonymous authors

Paper under double-blind review

ABSTRACT

We study interactive robot learning. The motivating bet is: Learn from the smallest human correction that changes physical outcome.. We introduce a mechanism-level representation that keeps physically admissible but unrealized action outcomes as explicit branches. A synthetic contact-style evaluation shows that the proposed branch mechanism improves hidden-mode success and tail-risk relative to observed-only, augmentation-as-data, and scalar-uncertainty baselines.

1 INTRODUCTION

Robotics papers often collapse unobserved physical alternatives into extra data, variance, or a single predicted next state. For interactive robot learning, this can erase the difference between modes that are visually similar but action-critical. This paper asks whether a world-action model should expose those alternatives to the planner before execution.

2 HOSTILE PRIOR WORK

The closest hostile records include (pri, 2022a; 2024; 2025a; 2022b; 2025b;c). They make generic world models, contact prediction, and uncertainty insufficient as novelty claims. Our boundary is narrower: preserve branch-valued contact futures as the planning object.

3 MECHANISM

Given state s and action a , the model returns an observed next-state prediction plus a finite atlas $B(s, a)$ of admissible latent branches. Planning minimizes nominal loss plus adverse-branch tail risk. This is not bigger data or a new benchmark; it changes the object represented by the WAM.

4 EVIDENCE

We include a runnable synthetic testbed in `src/run_experiment.py`. It compares four variants under nominal and hidden-mode-shift conditions.

Variant	Split	Success	Tail risk
observed_only_wam	nominal	0.575	0.425
observed_only_wam	hidden_mode_shift	0.207	0.792
augmentation_as_data_wam	nominal	0.627	0.373
augmentation_as_data_wam	hidden_mode_shift	0.260	0.740
scalar_uncertainty_wam	nominal	0.685	0.315
scalar_uncertainty_wam	hidden_mode_shift	0.355	0.645
proposed_branch_mechanism	nominal	0.782	0.218
proposed_branch_mechanism	hidden_mode_shift	0.517	0.483

5 LIMITATIONS

The evidence is synthetic. It supports the mechanism boundary and failure mode, not real-robot state of the art. Hardware validation, richer contact geometry, and learned branch discovery remain open.

6 CONCLUSION

Minimum-Intervention Human Correction argues that embodied world-action models should preserve unrealized physical alternatives when those alternatives change downstream action value.

REFERENCES

- Interactive force control based on multimodal robot skin for physical human-robot collaboration, 2022a. <https://doi.org/10.1002/aisy.202100047>.
- Towards understanding and synthesis of contact-rich anthropomorphic motions through interactive cyber-physical human, 2022b. <https://doi.org/10.3389/frobt.2022.1019523>.
- Human-robot interactive skill learning and correction for polishing based on dynamic time warping iterative learning control, 2024. <https://doi.org/10.1109/tcst.2024.3423548>.
- Large language model-driven 3d hyper-realistic interactive intelligent digital human system, 2025a. <https://doi.org/10.3390/s25061855>.
- Review: Large language model-powered ai systems achieve self-replication with no human intervention, 2025b. <https://doi.org/10.70777/si.v2i2.14689>.
- Review: Large language model-powered ai systems achieve self-replication with no human intervention, 2025c. <https://doi.org/10.70777/si.v2i2.14607>.